

Contextual Bandits Decision Tree:

Decision Node: ●

		Opens	Clicks	Mean
	A	4000	140	0.035
	B	3000	150	0.05

X	Option	Opens	Clicks	Mean
0	A	2500	50	0.02
	B	1000	100	0.10

X	Option	Opens	Clicks	Mean
1	A	1500	90	0.060
	B	2000	50	0.025

Strategy	Mean (μ)	Incremental Clicks
Action*	0.05	60
Action* X	0.08	210

Value Selection:

The expected reward (μ) when using variable X_i to select the best action is the weighted sum of expected rewards of the best actions for each of its possible values

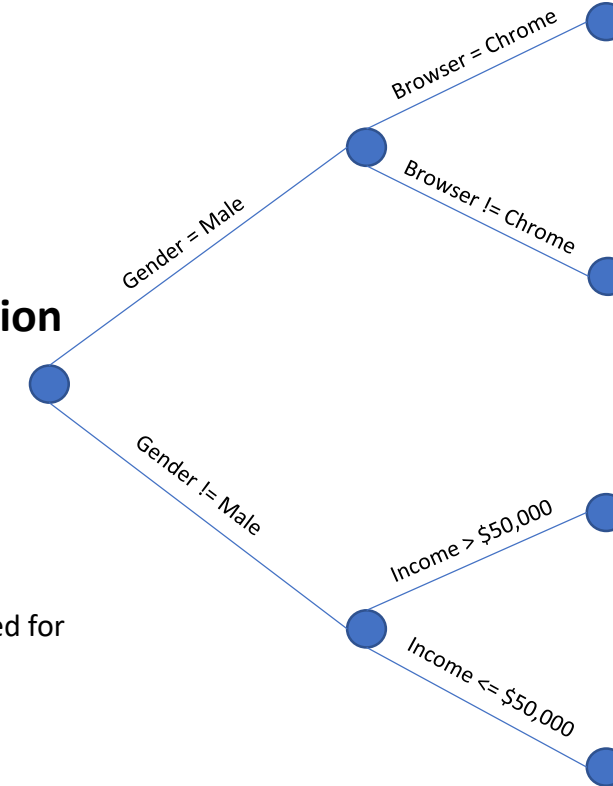
For continuous variables, a binary split is created to produce two segments of the data: $>$ or \leq value

For categorical variables, a one-vs-all strategy is used to create the two segments: $=$ or \neq value

Variable Selection:

The optimal variable to be used for selecting the best action is:
 $X_i^* = \text{argmax}(\mu | X_i)$

Decision Tree:







Thompson Sampling:

Each option is modeled by a Beta Distribution with:

$b = \text{\#opens}$

$a = \text{\#clicks}$



Original idea came from:
 Feraud, R. , Allesarido, R. , Urvoy, T. & Clerot, F.
Random Forest for the Contextual Bandit Problem